Contents

1   Unbiased Implicit Variational Inference

Based on Titsias and Ruiz [1].

- Authors introduce unbiased implicit variational inference (UIVI) that defines a flexible variational family. Like semi-implicit variational inference (SIVI), UIVI uses an implicit variational distribution $q_\theta(z) = \int q_\theta(z|\varepsilon)q(\varepsilon)d\varepsilon$ where $q_\theta(z|\varepsilon)$ is a reparameterizable distribution whose parameters can be outputs of some neural network g, i.e., $q_\theta(z|\varepsilon) = h(u; g(\varepsilon; \theta))$ with $u \sim q(u)$. Under two assumptions on the conditional $q_\theta(z|\varepsilon)$, the ELBO can be approximated via Monte Carlo sampling. In particular, the entropy component of the ELBO can be rewritten as an expectation w.r.t. the reverse conditional $q_\theta(\varepsilon|z)$. Efficient approximation of this expectation w.r.t. the reverse conditional is done by reusing samples from approximating the main expectation to initialize a MCMC sampler.

- Questions: TODO

    1. Can the gradient be pushed into the expectation? (Section 2.2)

- In SIVI, the variational distribution $q_\theta(z)$ is defined as

$$q_\theta(z) = \int q_\theta(z|\varepsilon)q(\varepsilon)d\varepsilon$$

where $\varepsilon \sim q(\varepsilon)$.

- UIVI:

    - Like SIVI, UIVI uses an implicit variational distribution $q_\theta(z)$ whose density cannot be evaluated but from which samples can be drawn. Unlike SIVI, UIVI directly maximizes the ELBO rather than a lower bound.

    - The dependence of $q_\theta(z|\varepsilon)$ on ε can be arbitrarily complex. Titsias and Ruiz [1] take the parameters of a reparameterizable distribution (Assumption 1) as the output of a neural network with parameters θ that takes ε as input, i.e.,

$$z = h(u; g_\theta(\varepsilon)) = h_\theta(u; \varepsilon)$$

where $u \sim q(u)$ and g_θ is some neural network. It is also assumed that $\nabla_z \log q_\theta(z|\varepsilon)$ can be evaluated (Assumption 2).

    - The gradient of the ELBO is given by

$$
\begin{aligned}
\nabla_\theta \mathcal{L}(\theta) &= \nabla_\theta \mathbb{E}_{q_\theta(z)}\left[\log p(x,z) - \log q_\theta(z)\right] \\
&= \nabla_\theta \mathbb{E}_{q(\varepsilon)q(u)}\left[\log p(x,z) - \log q_\theta(z)\big|_{z=h_\theta(u;\varepsilon)}\right] \\
&= \nabla_\theta \left(\mathbb{E}_{q(\varepsilon)q(u)}\left[\log p(x,z)\big|_{z=h_\theta(u;\varepsilon)}\right] - \mathbb{E}_{q(\varepsilon)q(u)}\left[\log q_\theta(z)\big|_{z=h_\theta(u;\varepsilon)}\right]\right) \\
&= \mathbb{E}_{q(\varepsilon)q(u)}\left[\nabla_z \log p(x,z)\big|_{z=h_\theta(u;\varepsilon)}\nabla_\theta h_\theta(u;\varepsilon)\right] - \mathbb{E}_{q(\varepsilon)q(u)}\left[\nabla_z \log q_\theta(z)\big|_{z=h_\theta(u;\varepsilon)}\nabla_\theta h_\theta(u;\varepsilon)\right] .
\end{aligned}
$$

(TODO: where is $\mathbb{E}_{q_\theta(z)}\left[\nabla_\theta \log q_\theta(z)\right] = 0$ applied?) (TODO: pushing gradient into expectation?) As $\nabla_z \log q_\theta(z)$ cannot be evaluated, this gradient is rewritten as an expectation using the log-

deritative identity: $\nabla_x \log f(x) = \frac{1}{f(x)} \nabla_x f(x)$:

$$
\begin{aligned}
\nabla_z \log q_\theta(z) &= \frac{1}{q_\theta(z)} \nabla_z q_\theta(z) \\
&= \frac{1}{q_\theta(z)} \nabla_z \int q_\theta(z|\varepsilon) q(\varepsilon) d\varepsilon \\
&= \frac{1}{q_\theta(z)} \int \nabla_z q_\theta(z|\varepsilon) q(\varepsilon) d\varepsilon \\
&= \frac{1}{q_\theta(z)} \int q_\theta(z|\varepsilon) q(\varepsilon) \nabla_z \log q_\theta(z|\varepsilon) d\varepsilon \\
&= \int q_\theta(\varepsilon|z) \nabla_z \log q_\theta(z|\varepsilon) d\varepsilon \\
&= \mathbb{E}_{q_\theta(\varepsilon|z)} \left[\nabla_z \log q_\theta(z|\varepsilon) \right] .
\end{aligned}
$$

$\nabla_z \log q_\theta(z|\varepsilon)$ can be evaluated by assumption.

- UIVI estimates the gradient of the ELBO by drawing S samples from $q(\varepsilon)$ and $q(u)$ (in practice, $S = 1$):

$$
\nabla_\theta \mathcal{L}(\theta) \approx \frac{1}{S} \sum_{s=1}^{S} \left(\nabla_z \log p(x, z) \big|_{z=h_\theta(u_s, \varepsilon_s)} \nabla_\theta h_\theta(u_s; \varepsilon_s) - \mathbb{E}_{q_\theta(\varepsilon|z)} \left[\nabla_z \log q_\theta(z|\varepsilon) \right] \big|_{z=h_\theta(u_s; \varepsilon_s)} \nabla_\theta h_\theta(u_s; \varepsilon_s) \right) .
$$

To estimate the inner expectation, samples are drawn from the reverse conditional $q_\theta(\varepsilon|z) \propto q_\theta(z|\varepsilon) q(\varepsilon)$ using MCMC. Exploiting the fact that (z_s, ε_s) comes from the joint $q_\theta(z, \varepsilon)$, UIVI initializes the MCMC at ε_s so no burn-in is required. A number of iterations are run to break the dependency between ε_s and the ε'_s that is used to estimate the inner expectation.

1.1 Analysis

TODO: analyze the (best-case) approximation of UIVI. Questions:

1. Approach? Probabilistic bound on KL as function of ELBO optimization iteration?

2. How to deal with implicit mixing component? Do surrogate families simpler than neural networks help? What assumptions would be needed?

2   Other references

VI review:

- Advances in Variational Inference (2019)

- Variational Inference: A Review for Statisticians (2017)

Possibly related VI approaches/of interest

- Semi-Implicit Variational Inference (2018)
  Doubly Semi-Implicit Variational Inference (2019)
  Structured Semi-Implicit Variational Inference (2019)
  Efficient Semi-Implicit Variational Inference (2021)

- Importance Weighted Hierarchical Variational Inference (2019)

- Stochastic Normalizing Flows (2020)

Theory/analysis

- Statistical Guarantees for Transformation Based Models with Applications to Implicit Variational Inference (2021)
  Statistical and Computational Properties of Variational Inference (2021; thesis)

- Theoretical Guarantees of Variational Inference and Its Applications (2020; thesis)

- Contributions to the theoretical study of variational inference and robustness (2020; thesis)

- On Statistical Optimality of Variational Bayes (2018)
  Statistical guarantees for variational Bayes (2021; slides)

- Statistical Guarantees and Algorithmic Convergence Issues of Variational Boosting (2020)

- Robust, Accurate Stochastic Optimization for Variational Inference (2020) – iterates as MCMC?

- Convergence Rates of Variational Inference in Sparse Deep Learning (2019)
  On the Convergence of Extended Variational Inference for Non-Gaussian Statistical Models (2020)

References

[1] Michalis K Titsias and Francisco Ruiz. Unbiased implicit variational inference. In *The 22nd International Conference on Artificial Intelligence and Statistics*, pages 167–176. PMLR, 2019.